

```
In [210]: import numpy as np
import pandas as pd
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import boxcox
```

```
In [44]: data=pd.read_csv('C:/Users/User/Downloads/G17MVSFTTRUCKS.csv',header=0,index_col=0)
```

```
In [45]: data.head()
```

```
Out[45]:
```

DATE	
1996-01-01	99.23
1996-02-01	105.93
1996-03-01	109.04
1996-04-01	106.96
1996-05-01	113.06

```
In [51]: len(data)
```

```
Out[51]: 330
```

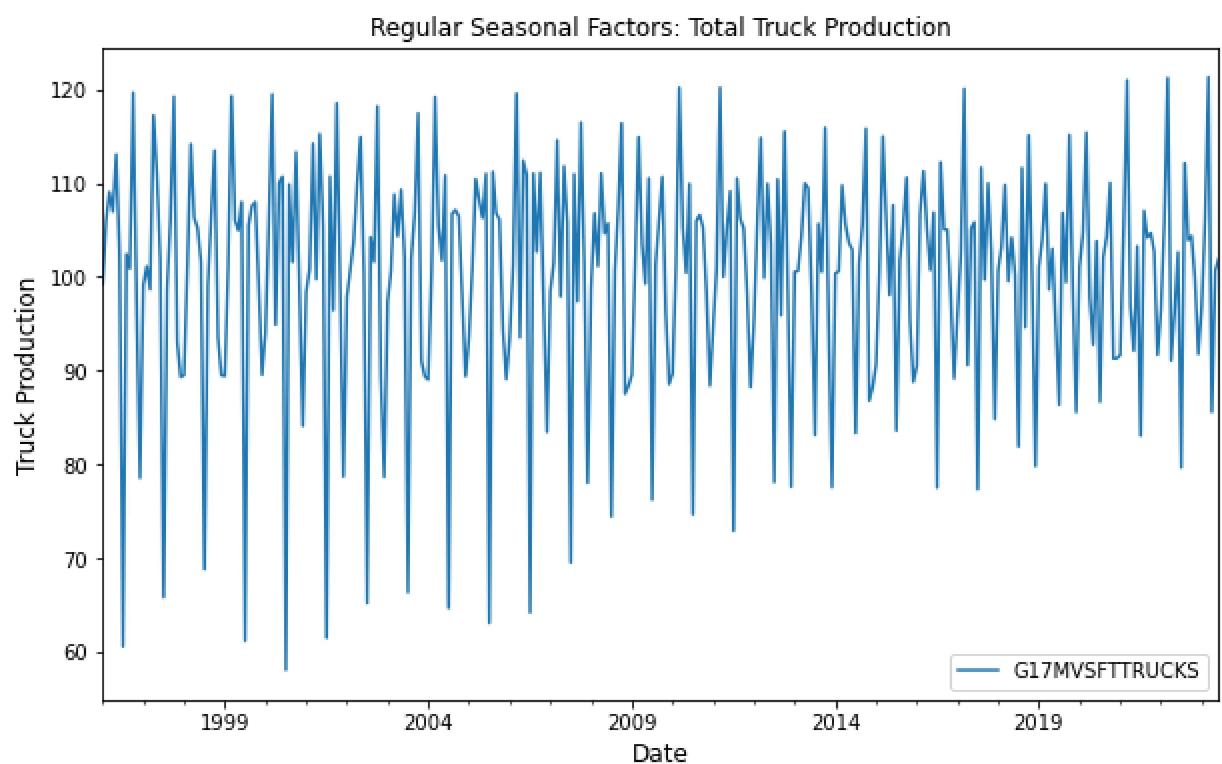
```
In [46]: data.describe()
```

```
Out[46]:
```

count	330.000000
mean	100.122758
std	12.062724
min	58.050000
25%	94.955000
50%	101.595000
75%	107.027500
max	121.320000

```
In [ ]: Time series plot
```

```
In [85]: data.plot()
plt.title('Regular Seasonal Factors: Total Truck Production')
plt.xlabel('Date', size = 12)
plt.ylabel('Truck Production', size = 12)
plt.rc("figure", figsize=(10,8))
```



```
In [ ]: We can observe from the above plot that the data is stationary.
```

Next we shall perform the Augmented Dickey Fuller test to ensure our data is stat

```
In [86]: from statsmodels.tsa.stattools import adfuller
```

```
In [87]: adfuller(data)
```

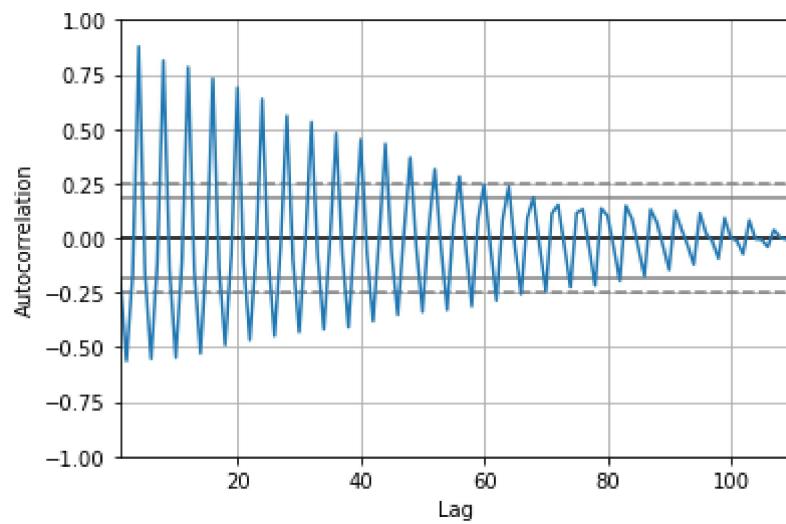
```
Out[87]: (-3.338469786748859,
 0.013244039555686926,
 17,
 312,
 {'1%': -3.4514843502727306,
 '5%': -2.8708485956333556,
 '10%': -2.571729625657462},
 1735.0980479552418)
```

```
In [ ]: The Augmented Dickey fuller test has a P value of greater less than 0.05 which se  
stationary.  
so we will move ahead to next step to plot the acf and pacf and find the models
```

```
In [89]: from statsmodels.graphics.tsaplots import plot_acf  
from statsmodels.graphics.tsaplots import plot_pacf
```

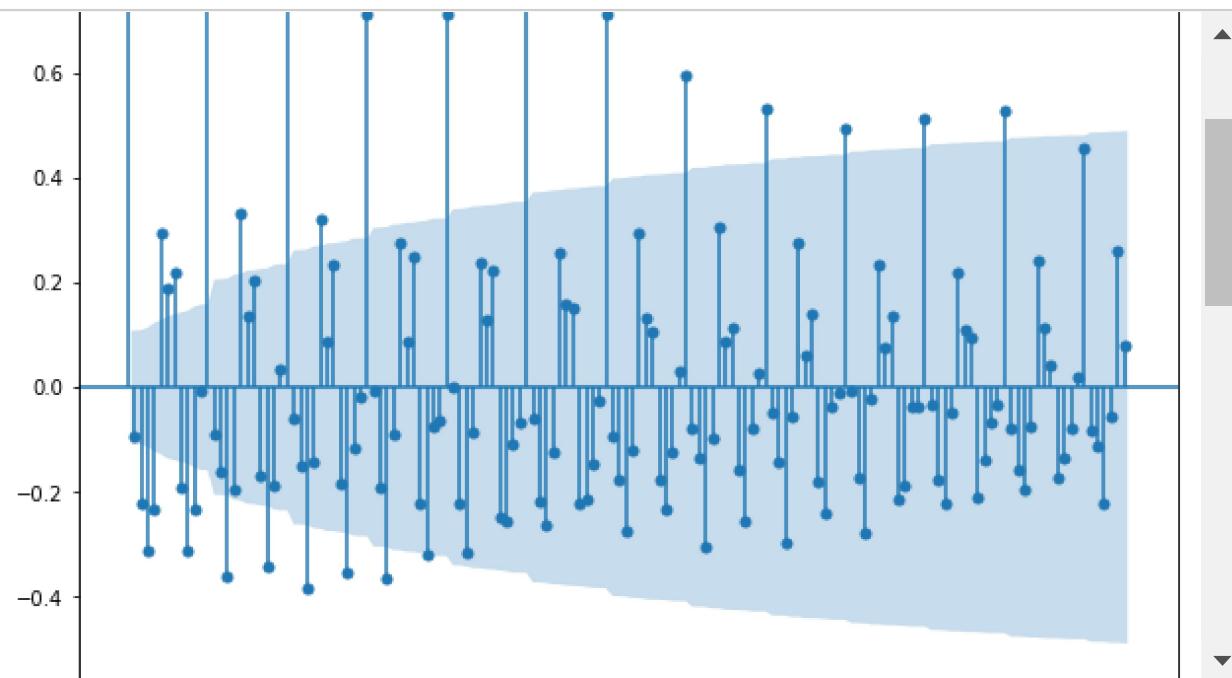
```
In [22]: import pandas  
from pandas.plotting import autocorrelation_plot  
from matplotlib import pyplot  
%matplotlib inline
```

```
In [23]: #Test  
autocorrelation_plot(data_q)  
pyplot.show()
```



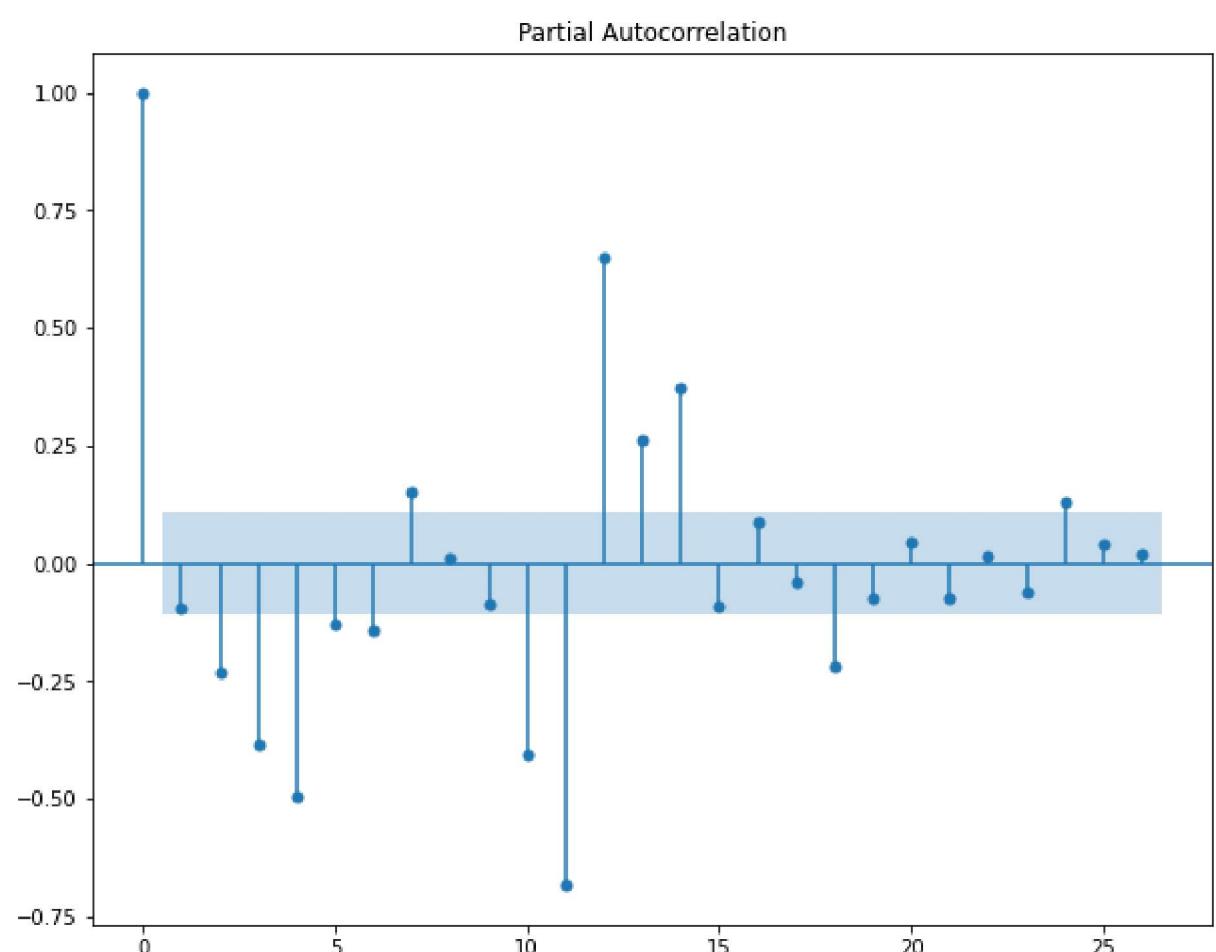
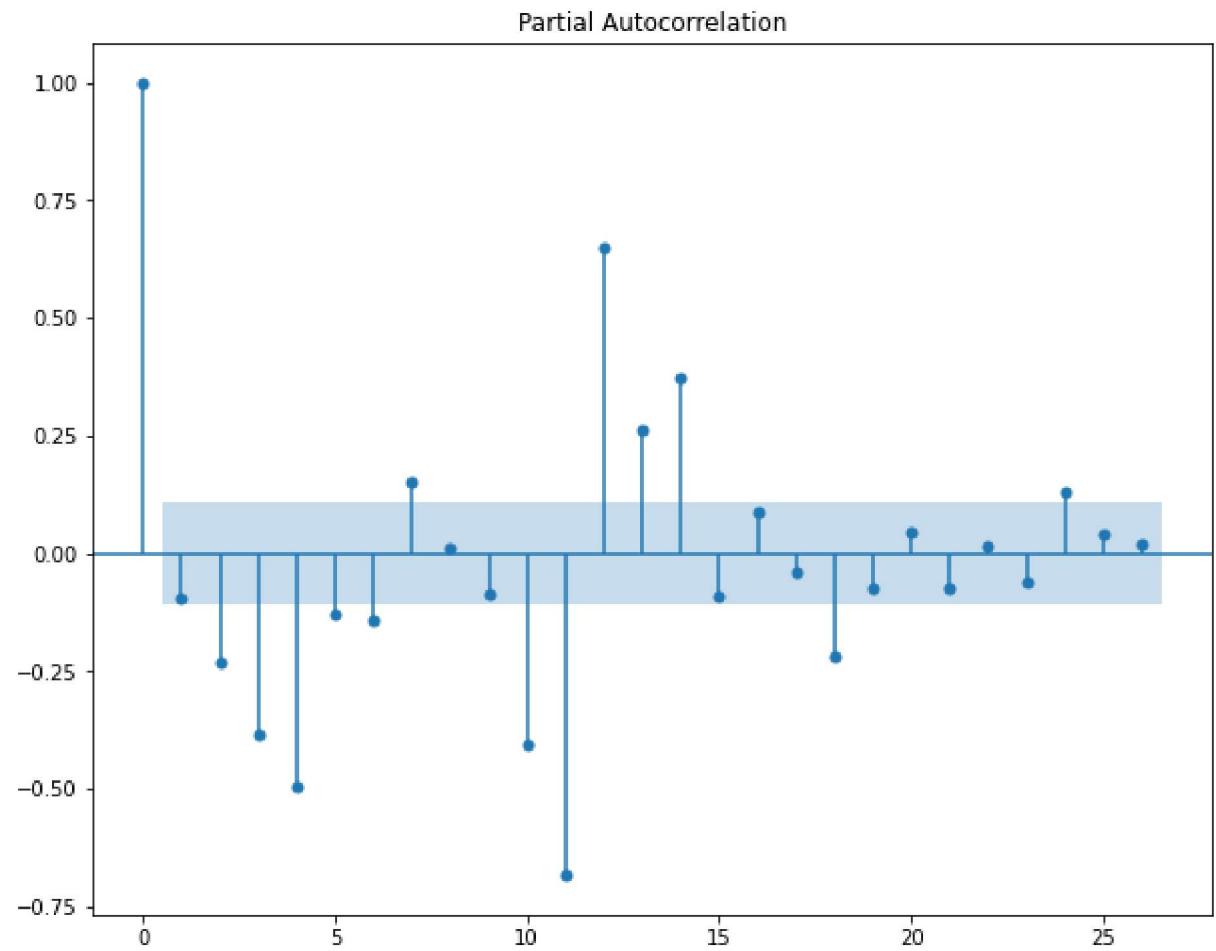
```
In [ ]: Plotting the ACF and Pacf
```

```
In [331]: plot_acf(data, lags=150)
```



In [94]: `plot_pacf(data)`

Out[94]:



```
In [ ]: Model Prediction by checking p,d,q values from ACF and PACF graph
```

```
In [96]: from pmдарима import auto_arima
import warnings
warnings.filterwarnings("ignore")
stepwise_fit=auto_arima(data['G17MVSFTTRUCKS'],trace=True,supress_warnings=True)
```

```
Performing stepwise search to minimize aic
ARIMA(2,0,2)(0,0,0)[0] intercept : AIC=2537.704, Time=0.57 sec
ARIMA(0,0,0)(0,0,0)[0] intercept : AIC=2582.977, Time=0.01 sec
ARIMA(1,0,0)(0,0,0)[0] intercept : AIC=2582.045, Time=0.03 sec
ARIMA(0,0,1)(0,0,0)[0] intercept : AIC=2546.157, Time=0.07 sec
ARIMA(0,0,0)(0,0,0)[0] intercept : AIC=3983.463, Time=0.01 sec
ARIMA(1,0,2)(0,0,0)[0] intercept : AIC=2549.563, Time=0.80 sec
ARIMA(2,0,1)(0,0,0)[0] intercept : AIC=2579.678, Time=0.26 sec
ARIMA(3,0,2)(0,0,0)[0] intercept : AIC=2567.545, Time=0.62 sec
ARIMA(2,0,3)(0,0,0)[0] intercept : AIC=2465.026, Time=0.70 sec
ARIMA(1,0,3)(0,0,0)[0] intercept : AIC=2489.930, Time=0.47 sec
ARIMA(3,0,3)(0,0,0)[0] intercept : AIC=2455.023, Time=0.72 sec
ARIMA(4,0,3)(0,0,0)[0] intercept : AIC=2409.031, Time=0.95 sec
ARIMA(4,0,2)(0,0,0)[0] intercept : AIC=2487.515, Time=0.73 sec
ARIMA(5,0,3)(0,0,0)[0] intercept : AIC=2361.373, Time=1.11 sec
ARIMA(5,0,2)(0,0,0)[0] intercept : AIC=2424.988, Time=1.02 sec
ARIMA(5,0,4)(0,0,0)[0] intercept : AIC=2439.187, Time=1.08 sec
ARIMA(4,0,4)(0,0,0)[0] intercept : AIC=2366.167, Time=0.97 sec
ARIMA(5,0,3)(0,0,0)[0] intercept : AIC=inf, Time=0.88 sec

Best model: ARIMA(5,0,3)(0,0,0)[0] intercept
Total fit time: 11.001 seconds
```

In [98]: `stepwise_fit.summary()`

Out[98]: SARIMAX Results

Dep. Variable:	y	No. Observations:	330			
Model:	SARIMAX(5, 0, 3)	Log Likelihood	-1170.687			
Date:	Mon, 16 May 2022	AIC	2361.373			
Time:	16:38:50	BIC	2399.364			
Sample:	0 - 330	HQIC	2376.527			
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
intercept	317.9807	33.258	9.561	0.000	252.796	383.165
ar.L1	-0.4832	0.108	-4.477	0.000	-0.695	-0.272
ar.L2	-0.1495	0.114	-1.306	0.192	-0.374	0.075
ar.L3	-0.9453	0.039	-23.946	0.000	-1.023	-0.868
ar.L4	-0.4472	0.091	-4.931	0.000	-0.625	-0.269
ar.L5	-0.1477	0.096	-1.544	0.123	-0.335	0.040
ma.L1	-0.0872	0.110	-0.796	0.426	-0.302	0.128
ma.L2	-0.1732	0.106	-1.632	0.103	-0.381	0.035
ma.L3	0.8566	0.094	9.108	0.000	0.672	1.041
sigma2	84.6304	9.433	8.971	0.000	66.141	103.119
Ljung-Box (L1) (Q):	0.01	Jarque-Bera (JB):	4.85			
Prob(Q):	0.90	Prob(JB):	0.09			
Heteroskedasticity (H):	0.43	Skew:	0.24			
Prob(H) (two-sided):	0.00	Kurtosis:	3.35			

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [412]: model=ARIMA(data['G17MVSFTTRUCKS'],order=(4,0,4))
model=model.fit()
model.summary()
```

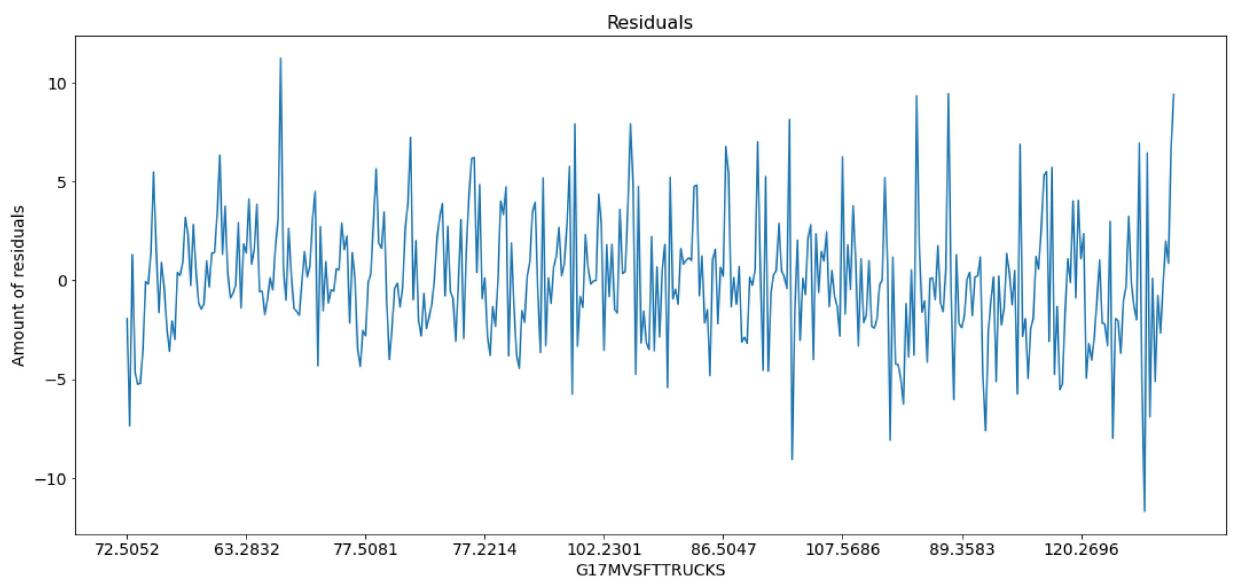
Out[412]: SARIMAX Results

Dep. Variable:		G17MVSFTTRUCKS	No. Observations:	330		
Model:		ARIMA(4, 0, 4)	Log Likelihood	-1135.360		
Date:		Mon, 16 May 2022	AIC	2290.720		
Time:		23:12:20	BIC	2328.711		
Sample:		01-01-1996 - 06-01-2023	HQIC	2305.874		
Covariance Type:		opg				
	coef	std err	z	P> z	[0.025	0.975]
const	100.0955	0.333	300.652	0.000	99.443	100.748
ar.L1	-0.7377	0.013	-55.287	0.000	-0.764	-0.712
ar.L2	-0.2784	0.022	-12.560	0.000	-0.322	-0.235
ar.L3	-0.7402	0.013	-57.224	0.000	-0.766	-0.715
ar.L4	-0.9966	0.003	-332.261	0.000	-1.002	-0.991
ma.L1	0.5217	48.691	0.011	0.991	-94.911	95.955
ma.L2	-0.1375	5.165	-0.027	0.979	-10.261	9.986
ma.L3	0.5314	90.666	0.006	0.995	-177.170	178.233
ma.L4	0.9921	79.088	0.013	0.990	-154.018	156.002
sigma2	52.2048	4161.318	0.013	0.990	-8103.829	8208.238
Ljung-Box (L1) (Q):		9.11	Jarque-Bera (JB):	4.27		
Prob(Q):		0.00	Prob(JB):	0.12		
Heteroskedasticity (H):		0.46	Skew:	-0.28		
Prob(H) (two-sided):		0.00	Kurtosis:	2.98		

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [425]: # Plotting residual
residuals = pd.DataFrame(model2_fit.resid)
plt.figure(figsize=(18, 8))
plt.plot(residuals)
plt.title('Residuals', fontsize=16)
plt.xlabel("G17MVSFTTRUCKS", fontsize=14)
plt.ylabel("Amount of residuals", fontsize=14)
plt.xticks(np.arange(0, len(data2.IPG2211A2N)+1, 45), labels=[data2.IPG2211A2N[i] for i in range(0, len(data2.IPG2211A2N)+1)])
plt.yticks(fontsize=14)
plt.show()
```



```
In [99]: print(data.shape)
train= data.iloc[:,-30]
test= data.iloc[-30:]
print(train.shape,test.shape)
```

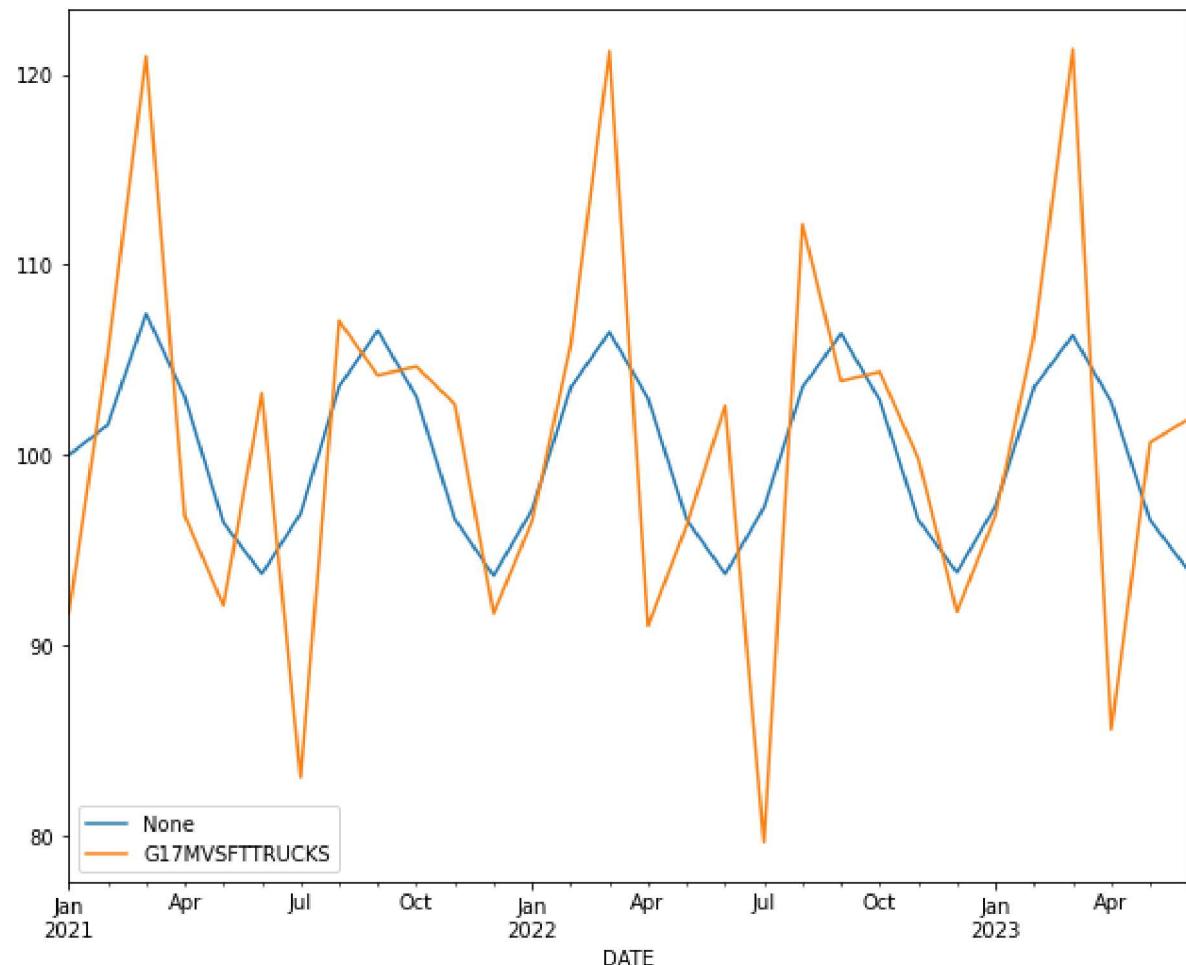
```
(330, 1)
(300, 1) (30, 1)
```

```
In [ ]: model=ARIMA(train[ 'G17MVSFTTRUCKS' ],order=(5,0,3))
model=model.fit()
model.summary()
```

```
In [370]: start=len(train)
end=len(train)+len(test)
predictions = model.predict(start=start, end=end, dynamic=False, typ='levels')
```

```
In [102]: pred.plot(legend=True)
test['G17MVSFTTRUCKS'].plot(legend=True)
```

```
Out[102]: <AxesSubplot:xlabel='DATE'>
```



```
In [ ]: From the above graph we can say that the actual and predicted values looks almost
Hence we can say it is the best model
```

```
In [39]: #Non-Seasonal Dataset
```

```
In [227]: data2=pd.read_csv('C:/Users/User/Downloads/Electric_Production.csv',header=0,pars  
print(data2)
```

	DATE	IPG2211A2N
0	1/1/1985	72.5052
1	2/1/1985	70.6720
2	3/1/1985	62.4502
3	4/1/1985	57.4714
4	5/1/1985	55.3151
..	...	...
392	9/1/2017	98.6154
393	10/1/2017	93.6137
394	11/1/2017	97.3359
395	12/1/2017	114.7212
396	1/1/2018	129.4048

[397 rows × 2 columns]

```
In [423]: data2=pd.read_csv('C:/Users/User/Downloads/Electric_Production.csv',header=0,inde  
data2
```

Out[423]:

IPG2211A2N

DATE	IPG2211A2N
1985-01-01	72.5052
1985-02-01	70.6720
1985-03-01	62.4502
1985-04-01	57.4714
1985-05-01	55.3151
...	...
2017-09-01	98.6154
2017-10-01	93.6137
2017-11-01	97.3359
2017-12-01	114.7212
2018-01-01	129.4048

397 rows × 1 columns

```
In [128]: data2.dtypes=='object'  
num_vars=data2.columns[data2.dtypes!='object']  
data2[num_vars]
```

Out[128]:

IPG2211A2N

DATE	IPG2211A2N
1985-01-01	72.5052
1985-02-01	70.6720
1985-03-01	62.4502
1985-04-01	57.4714
1985-05-01	55.3151
...	...
2017-09-01	98.6154
2017-10-01	93.6137
2017-11-01	97.3359
2017-12-01	114.7212
2018-01-01	129.4048

397 rows × 1 columns

```
In [129]: #Checking the null values in our dataset  
data2[num_vars].isnull()
```

Out[129]:

IPG2211A2N

DATE	
1985-01-01	False
1985-02-01	False
1985-03-01	False
1985-04-01	False
1985-05-01	False
...	...
2017-09-01	False
2017-10-01	False
2017-11-01	False
2017-12-01	False
2018-01-01	False

397 rows × 1 columns

```
In [130]: data2.isnull().sum()
```

Out[130]: IPG2211A2N 0  
dtype: int64

```
In [131]: data2.head()
```

IPG2211A2N

DATE	
1985-01-01	72.5052
1985-02-01	70.6720
1985-03-01	62.4502
1985-04-01	57.4714
1985-05-01	55.3151

```
In [132]: data2.describe()
```

Out[132]:

IPG2211A2N

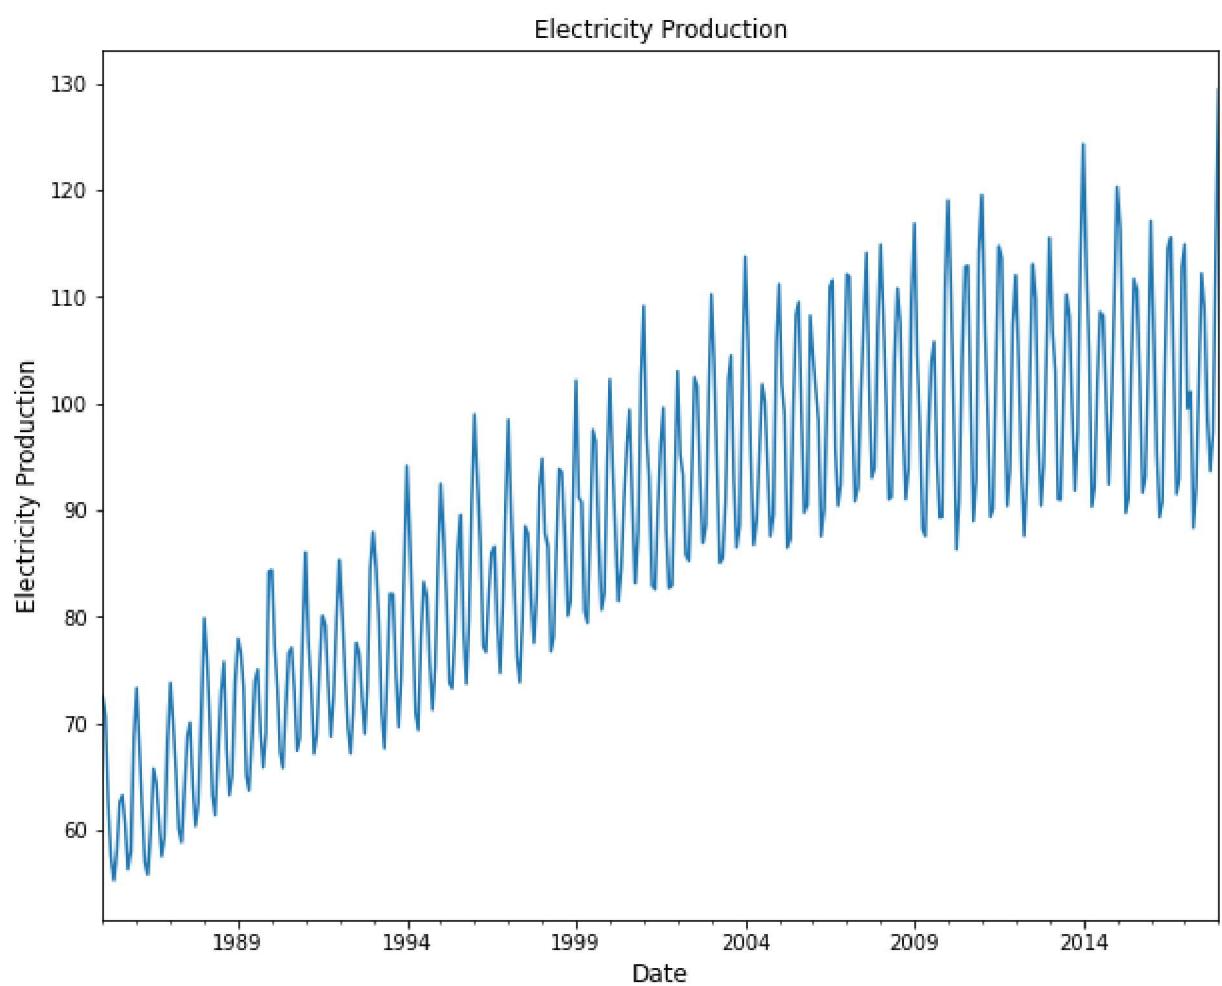
count	397.000000
mean	88.847218
std	15.387834
min	55.315100
25%	77.105200
50%	89.779500
75%	100.524400
max	129.404800

```
In [133]: data2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 397 entries, 1985-01-01 to 2018-01-01
Data columns (total 1 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   IPG2211A2N  397 non-null    float64
dtypes: float64(1)
memory usage: 6.2 KB
```

```
In [ ]: Time series plot
```

```
In [424]: y=data2.IPG2211A2N
y.plot()
plt.rc("figure", figsize=(10,8))
plt.title('Electricity Production')
plt.xlabel('Date', size = 12)
plt.ylabel('Electricity Production', size = 12)
plt.rc("figure", figsize=(10,8))
```



```
In [ ]: Looking at the plot we can observe there is an **upward trend** over the period of time.
```

#### Stationarity Test

We can observe **from** the above plot that the Electric production **is** fairly seasonal. Next we shall perform the Augmented Dickey Fuller test to see **if** the trend of the **is** stationary **or** non stationary.

```
In [ ]: The augmented Dickey-Fuller (ADF) test statistic is the t-statistic of the estimated least squares regression. However, the ADF test statistic is not approximately t-distributed under the null hypothesis; instead, it has a certain nonstandard large-sample distribution under the null hypothesis
```

```
In [229]: from statsmodels.tsa.stattools import adfuller
```

```
In [293]: data2=pd.read_csv('C:/Users/User/Downloads/Electric_Production.csv',header=0,inde  
adfuller(data2)
```

```
Out[293]: (-2.2569903500472455,  
 0.18621469116586759,  
 15,  
 381,  
 {'1%': -3.4476305904172904,  
 '5%': -2.869155980820355,  
 '10%': -2.570827146203181},  
 1840.8474501627156)
```

In [ ]: The Augmented Dickey fuller test has a P value of greater than **0.05** which seems to indicate that the time series **is non-stationary**. The p-value **is** obtained **is** greater than significance level of **any** of the critical values.

Clearly, there **is** no reason to reject the null hypothesis. So, the time series **is non-stationary**.

We can clearly see a trend **in** data so let us perform some more formal test of stationarity.

```
In [291]: #first differencing  
first_difference=adfuller(data2.IPG2211A2N.diff().dropna())  
print(f'ADF Statistics:{result[0]}')  
print(f'p-value:{result[1]}')
```

```
ADF Statistics:-6.748333370019157  
p-value:2.9951614981156204e-09
```

In [ ]: The p-value:**2.9951614981156204e-09** **is** less than **0.05**. Now the series looks stationary so q will be **1 for** our model. I am also testing another method two make data stationary.

```
In [236]: #Method2 only for Test  
#Differencing and then log transformation  
#Making data Stationary using Differencing  
from statsmodels.tsa.stattools import adfuller  
y1=y-y.shift(1)  
y1=y1[1:]  
y2=np.log(y1-np.min(y1)+1)
```

```
In [237]: result=adfuller(y1)  
print('ADF statistic: %f' %result[0])  
print('p-value: %f' %result[1])  
print('critical values:')
```

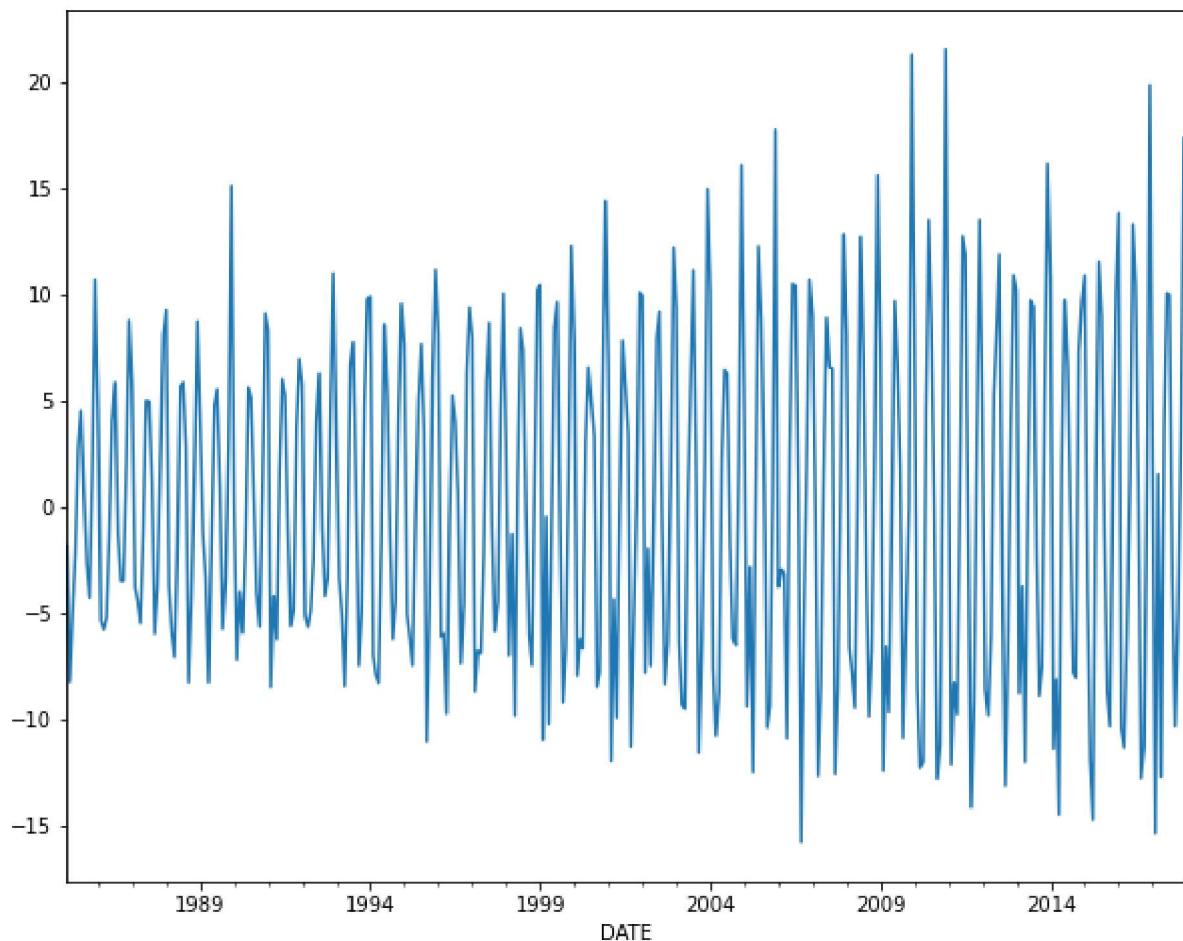
```
for key,value in result[4].items():  
    print('\t%s: %.3f' % (key,value))
```

```
ADF statistic: -7.104891  
p-value: 0.000000  
critical values:  
    1%: -3.448  
    5%: -2.869  
    10%: -2.571
```

```
In [238]: y1.plot()
```

```
Out[238]: <AxesSubplot:xlabel='DATE'>
```



```
In [ ]: The data looks stationary now but we can also try First Log Transformation  
and then differencing (only for test purpose and I am taking d as 1 as data change  
in first differencing)
```

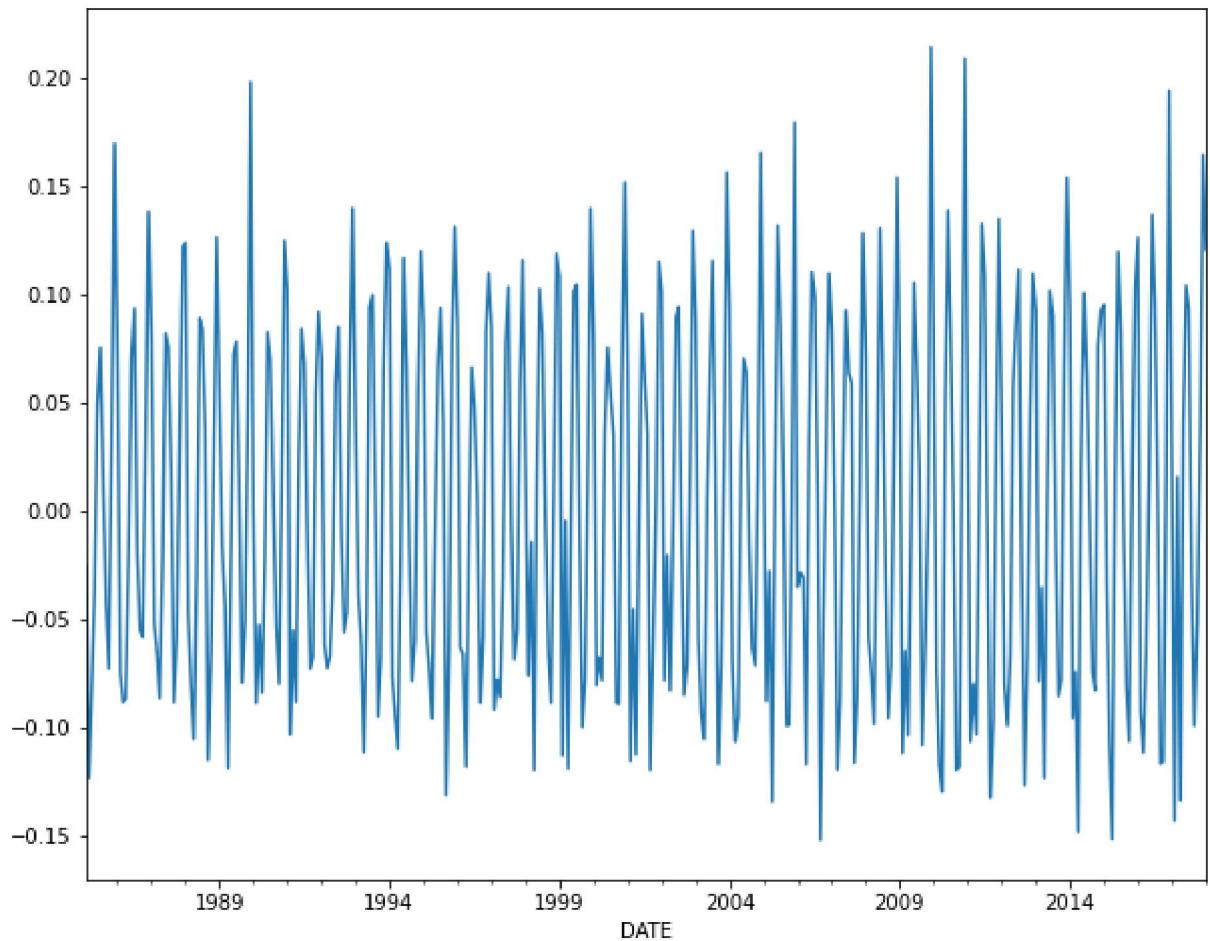
```
In [239]: from statsmodels.tsa.stattools import adfuller  
y1=np.log(y)  
y2=y1-y1.shift(1)  
y2.dropna(inplace=True)
```

```
In [240]: result=adfuller(y2)  
  
print('ADF statistic: %f' %result[0])  
print('p-value: %f' %result[1])  
print('critical values: ')  
  
for key,value in result[4].items():  
    print('\t%s: %.3f' % (key,value))
```

```
ADF statistic: -6.748333  
p-value: 0.000000  
critical values:  
    1%: -3.448  
    5%: -2.869  
    10%: -2.571
```

```
In [241]: y2.plot()
```

```
Out[241]: <AxesSubplot:xlabel='DATE'>
```



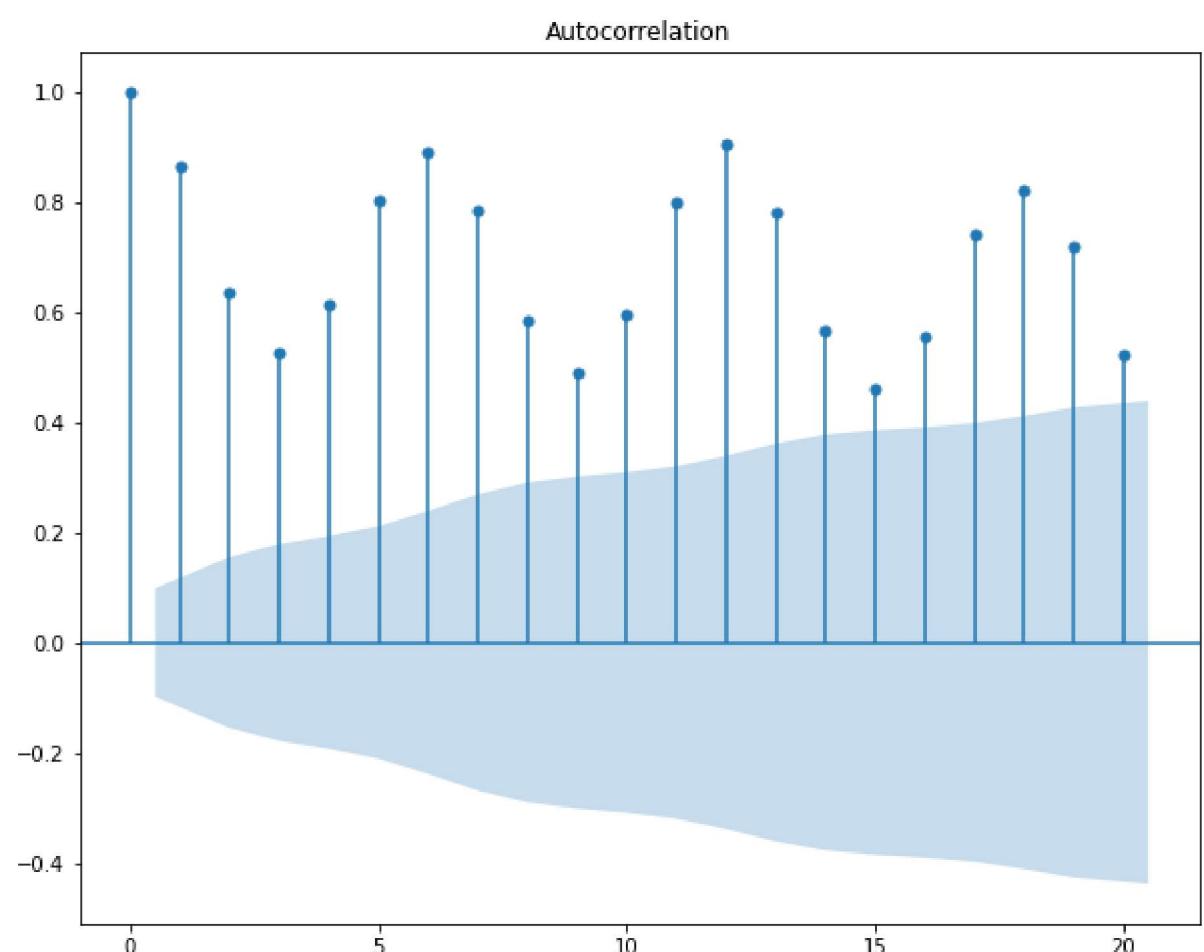
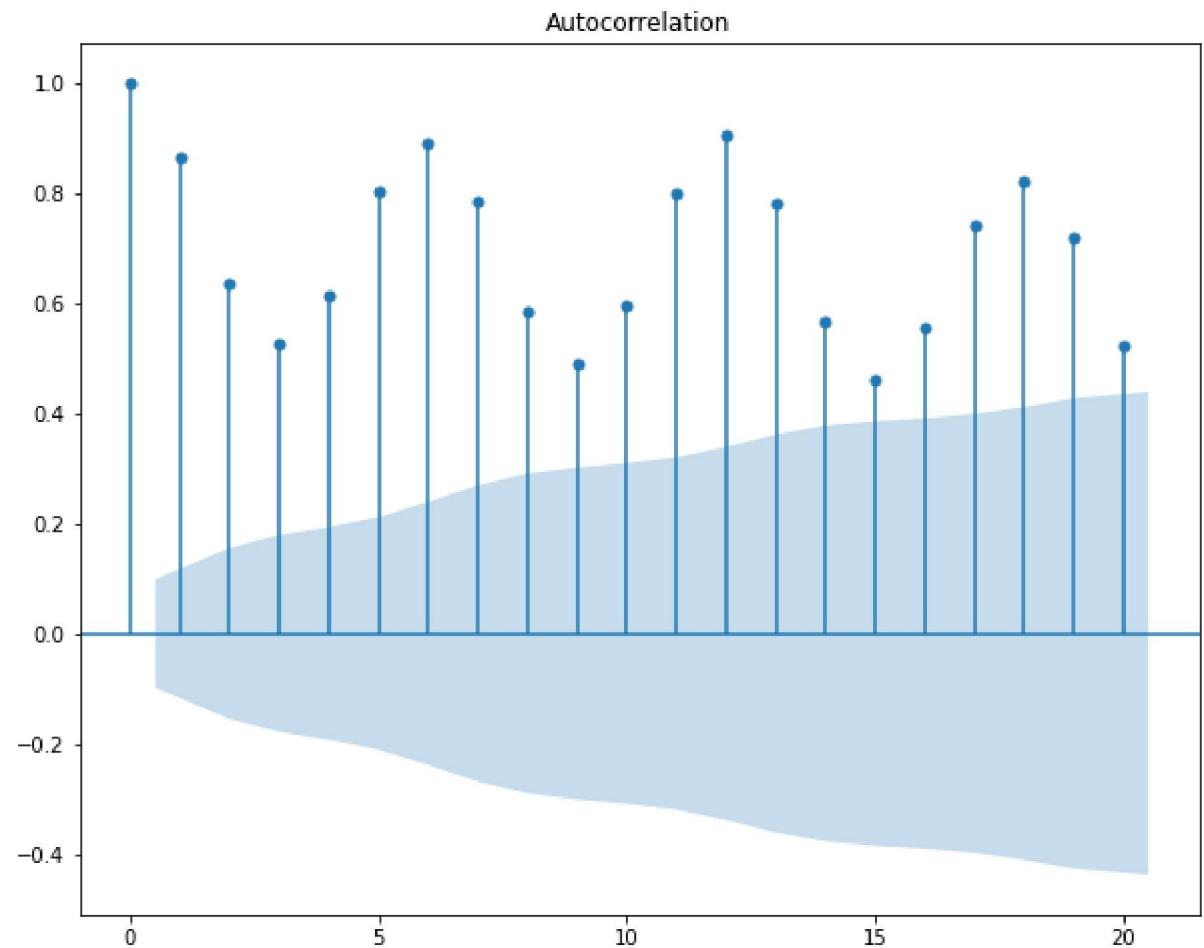
```
In [142]: print("Data Shape: {}".format(data2.shape))
value_1 = data2[0:198]
value_2 = data2[199:398]
```

```
Data Shape: (397, 1)
```

```
In [242]: from statsmodels.graphics.tsaplots import plot_acf
from statsmodels.graphics.tsaplots import plot_pacf
```

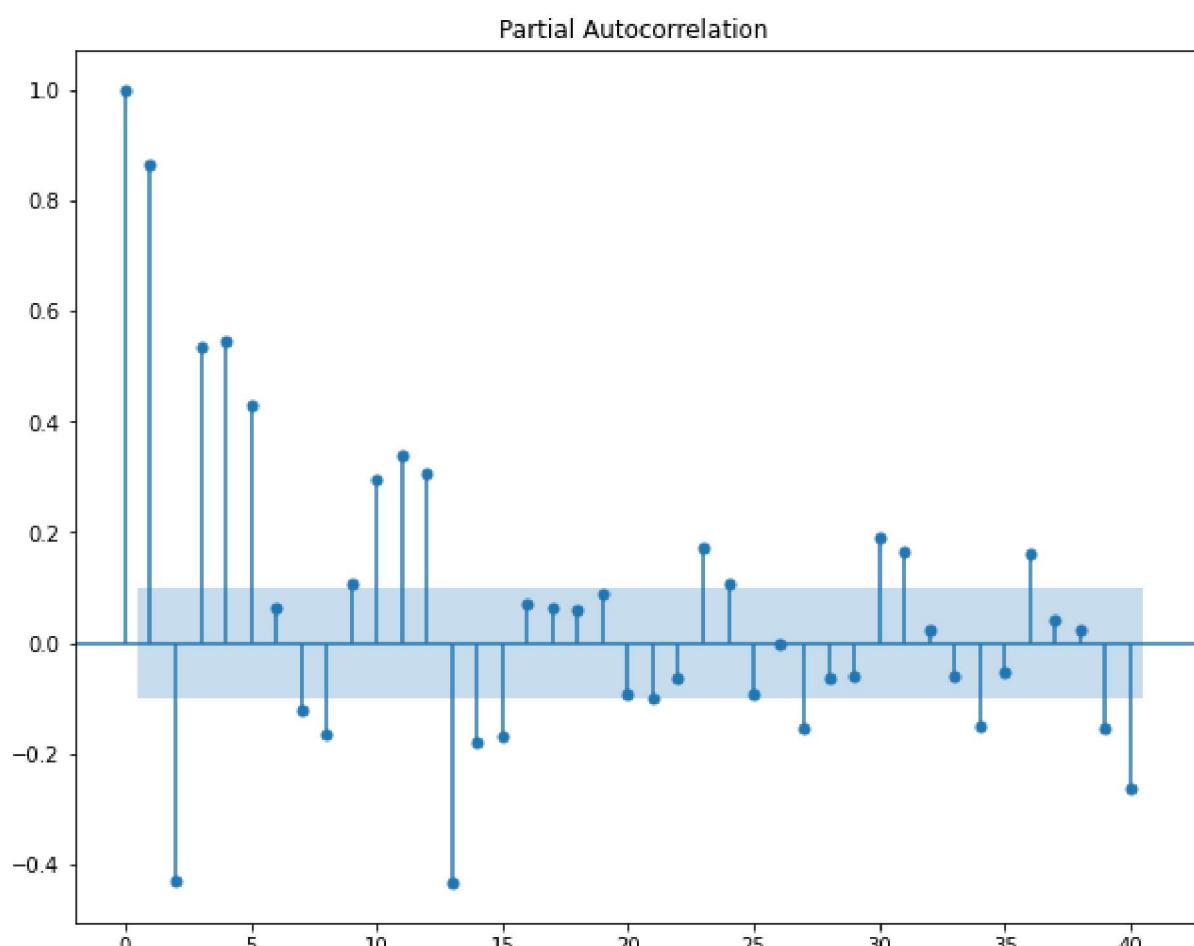
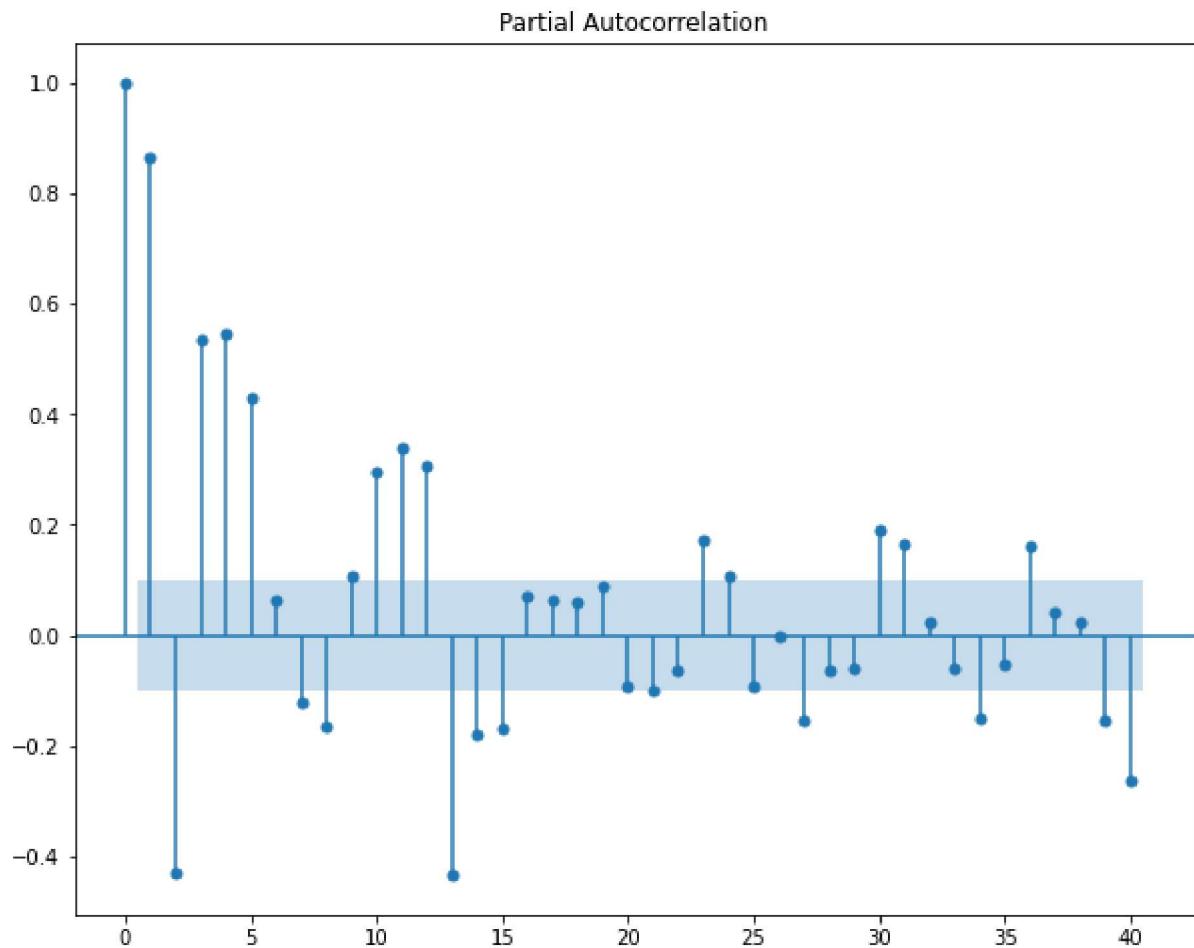
```
In [247]: plot_acf(data2, lags=20)
```

Out[247]:



```
In [246]: plot_pacf(data2, lags=40)
```

Out[246]:



```
In [ ]: Forecasting models
```

```
"Pmdarima" is a useful package to help us find the best parameters for SARIMA mod
```

```
In [157]: !pip install pmdarima
```

```
Requirement already satisfied: pmdarima in c:\users\user\anaconda3\lib\site-packages (1.8.5)
Requirement already satisfied: scipy>=1.3.2 in c:\users\user\anaconda3\lib\site-packages (from pmdarima) (1.6.2)
Requirement already satisfied: Cython!=0.29.18,>=0.29 in c:\users\user\anaconda3\lib\site-packages (from pmdarima) (0.29.23)
Requirement already satisfied: urllib3 in c:\users\user\anaconda3\lib\site-packages (from pmdarima) (1.26.4)
Requirement already satisfied: numpy>=1.19.3 in c:\users\user\anaconda3\lib\site-packages (from pmdarima) (1.20.1)
Requirement already satisfied: pandas>=0.19 in c:\users\user\anaconda3\lib\site-packages (from pmdarima) (1.2.4)
Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in c:\users\user\anaconda3\lib\site-packages (from pmdarima) (52.0.0.post20210125)
Requirement already satisfied: scikit-learn>=0.22 in c:\users\user\anaconda3\lib\site-packages (from pmdarima) (0.24.1)
Requirement already satisfied: joblib>=0.11 in c:\users\user\anaconda3\lib\site-packages (from pmdarima) (1.0.1)
Requirement already satisfied: statsmodels!=0.12.0,>=0.11 in c:\users\user\anaconda3\lib\site-packages (from pmdarima) (0.12.2)
Requirement already satisfied: pytz>=2017.3 in c:\users\user\anaconda3\lib\site-packages (from pandas>=0.19->pmdarima) (2021.1)
Requirement already satisfied: python-dateutil>=2.7.3 in c:\users\user\anaconda3\lib\site-packages (from pandas>=0.19->pmdarima) (2.8.1)
Requirement already satisfied: six>=1.5 in c:\users\user\anaconda3\lib\site-packages (from python-dateutil>=2.7.3->pandas>=0.19->pmdarima) (1.15.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\user\anaconda3\lib\site-packages (from scikit-learn>=0.22->pmdarima) (2.1.0)
Requirement already satisfied: patsy>=0.5 in c:\users\user\anaconda3\lib\site-packages (from statsmodels!=0.12.0,>=0.11->pmdarima) (0.5.1)
```

```
In [158]: from pmdarima import auto_arima
import warnings
warnings.filterwarnings("ignore")
```

```
In [ ]: ARIMA Forecast to get best p,d,q,P,D,Q values
```

```
In [253]: from statsmodels.tsa.arima_model import ARIMA
from statsmodels.tsa.statespace.sarimax import SARIMAX
```

```
In [283]: from statsmodels.tsa.arima_model import ARIMA
model1=ARIMA(data2['IPG2211A2N'],order=(2,1,4))
model1_fit=model1.fit()
print(model1_fit.summary())
```

ARIMA Model Results						
Dep. Variable:	D.IPG2211A2N	No. Observations:	396			
Model:	ARIMA(2, 1, 4)	Log Likelihood	-1032.303			
Method:	css-mle	S.D. of innovations	11550678.802			
Date:	Mon, 16 May 2022	AIC	2080.605			
Time:	21:08:31	BIC	2112.457			
Sample:	02-01-1985 - 01-01-2018	HQIC	2093.224			
	coef	std err	z	P> z	[0.025	
0.975]						
-----						
const	0.1075	4.35e-05	2472.383	0.000	0.107	
0.108						
ar.L1.D.IPG2211A2N	0.9997	0.000	2472.306	0.000	0.999	
1.000						
ar.L2.D.IPG2211A2N	-1.0000	nan	nan	nan	nan	
nan						
ma.L1.D.IPG2211A2N	-1.3643	0.009	-151.276	0.000	-1.382	
-1.347						
ma.L2.D.IPG2211A2N	0.9130	nan	nan	nan	nan	
nan						
ma.L3.D.IPG2211A2N	0.0126	0.002	5.872	0.000	0.008	
0.017						
ma.L4.D.IPG2211A2N	-0.4326	0.008	-51.455	0.000	-0.449	
-0.416						
	Roots					
-----						
	Real	Imaginary	Modulus	Frequency		
-----						
AR.1	0.4998	-0.8661j	1.0000	-0.1667		
AR.2	0.4998	+0.8661j	1.0000	0.1667		
MA.1	1.0920	-0.0000j	1.0920	-0.0000		
MA.2	0.4877	-0.8948j	1.0191	-0.1706		
MA.3	0.4877	+0.8948j	1.0191	0.1706		
MA.4	-2.0383	-0.0000j	2.0383	-0.5000		
-----						

In [ ]: When we look at the plot we can see there **is** a seasonality **in** data. That **is** why we used **Seasonal** instead of ARIMA.

There are four seasonal elements that are **not** part of ARIMA that must be configured:  
P: Seasonal autoregressive order.  
D: Seasonal difference order.  
Q: Seasonal moving average order.  
m: The number of time steps **for** a single seasonal period.

Now we will **try** to get best p,d,q,P,D,Q values based on the acf **and** pacf plots

```
In [284]: #ARIMA(1,1,1)(3,1,1)[6]
model2=SARIMAX(data2,order=(1,1,1), seasonal_order=(3,1,1,6))
results=model.fit()
results.summary()
```

Out[284]: SARIMAX Results

Dep. Variable:	IPG2211A2N	No. Observations:	397			
Model:	SARIMAX(1, 1, 1)x(3, 1, 1, 6)	Log Likelihood	-925.755			
Date:	Mon, 16 May 2022	AIC	1865.510			
Time:	21:09:04	BIC	1893.273			
Sample:	01-01-1985 - 01-01-2018	HQIC	1876.515			
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.5433	0.041	13.167	0.000	0.462	0.624
ma.L1	-0.9645	0.015	-62.294	0.000	-0.995	-0.934
ar.S.L6	-0.2686	0.063	-4.252	0.000	-0.392	-0.145
ar.S.L12	0.2281	0.064	3.568	0.000	0.103	0.353
ar.S.L18	-0.2357	0.051	-4.618	0.000	-0.336	-0.136
ma.S.L6	-0.8015	0.050	-16.004	0.000	-0.900	-0.703
sigma2	6.4779	0.376	17.211	0.000	5.740	7.216
Ljung-Box (L1) (Q):	0.27	Jarque-Bera (JB):	36.75			
Prob(Q):	0.60	Prob(JB):	0.00			
Heteroskedasticity (H):	2.75	Skew:	0.24			
Prob(H) (two-sided):	0.00	Kurtosis:	4.43			

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

In [ ]: The model summary provides lot of information. The table **in the middle** **is** the coefficients under ‘coef’ are the weights of the respective terms.

The P-Value **in** ‘P>|z|’ column **is** highly significant **as** it should be less than **0.05**.

So, we will rebuild the model.

```
In [287]: #ARIMAX(4, 1, 4)x(4, 0, [1], 12)
model2=SARIMAX(data2,order=(4,1,4), seasonal_order=(4,0,1,12))
results=model.fit()
results.summary()
```

Out[287]: SARIMAX Results

Dep. Variable:	IPG2211A2N	No. Observations:	397			
Model:	SARIMAX(1, 1, 1)x(3, 1, 1, 6)	Log Likelihood	-925.755			
Date:	Mon, 16 May 2022	AIC	1865.510			
Time:	21:16:28	BIC	1893.273			
Sample:	01-01-1985 - 01-01-2018	HQIC	1876.515			
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.5433	0.041	13.167	0.000	0.462	0.624
ma.L1	-0.9645	0.015	-62.294	0.000	-0.995	-0.934
ar.S.L6	-0.2686	0.063	-4.252	0.000	-0.392	-0.145
ar.S.L12	0.2281	0.064	3.568	0.000	0.103	0.353
ar.S.L18	-0.2357	0.051	-4.618	0.000	-0.336	-0.136
ma.S.L6	-0.8015	0.050	-16.004	0.000	-0.900	-0.703
sigma2	6.4779	0.376	17.211	0.000	5.740	7.216

Ljung-Box (L1) (Q): 0.27 Jarque-Bera (JB): 36.75  
Prob(Q): 0.60 Prob(JB): 0.00  
Heteroskedasticity (H): 2.75 Skew: 0.24  
Prob(H) (two-sided): 0.00 Kurtosis: 4.43

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [326]: model=SARIMAX(data2,order=(2,1,2), seasonal_order=(1, 1, 2, 12))
results=model.fit()
```

```
In [327]: results.summary()
```

```
Out[327]: SARIMAX Results
```

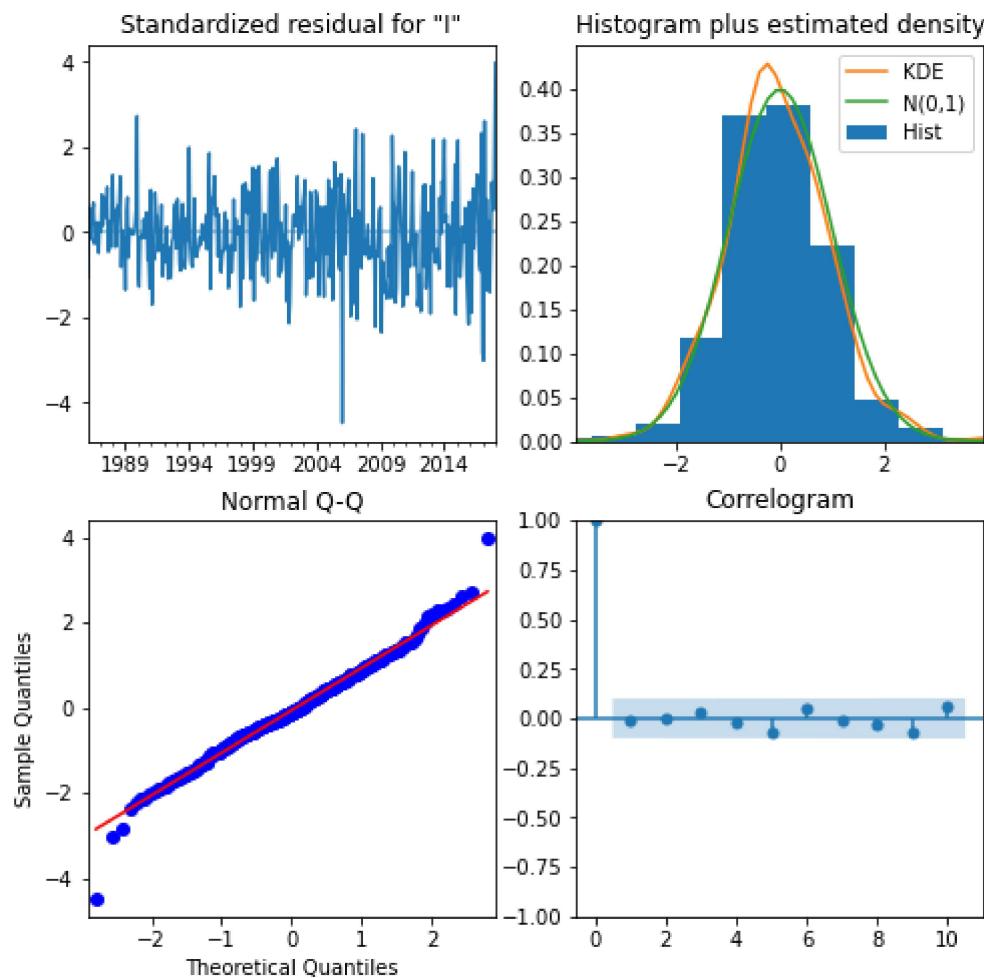
Dep. Variable:	IPG2211A2N	No. Observations:	397			
Model:	SARIMAX(2, 1, 2)x(1, 1, 2, 12)	Log Likelihood	-884.922			
Date:	Mon, 16 May 2022	AIC	1785.844			
Time:	22:14:35	BIC	1817.449			
Sample:	01-01-1985 - 01-01-2018	HQIC	1798.380			
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.0383	0.407	0.094	0.925	-0.759	0.836
ar.L2	0.1850	0.241	0.768	0.443	-0.287	0.657
ma.L1	-0.4249	0.398	-1.069	0.285	-1.204	0.354
ma.L2	-0.4995	0.385	-1.299	0.194	-1.253	0.254
ar.S.L12	-0.4947	0.334	-1.482	0.138	-1.149	0.160
ma.S.L12	-0.1772	0.331	-0.536	0.592	-0.826	0.471
ma.S.L24	-0.4653	0.238	-1.959	0.050	-0.931	0.000
sigma2	5.6905	0.334	17.057	0.000	5.037	6.344
Ljung-Box (L1) (Q):	0.06	Jarque-Bera (JB):	28.47			
Prob(Q):	0.81	Prob(JB):	0.00			
Heteroskedasticity (H):	3.01	Skew:	0.00			
Prob(H) (two-sided):	0.00	Kurtosis:	4.33			

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

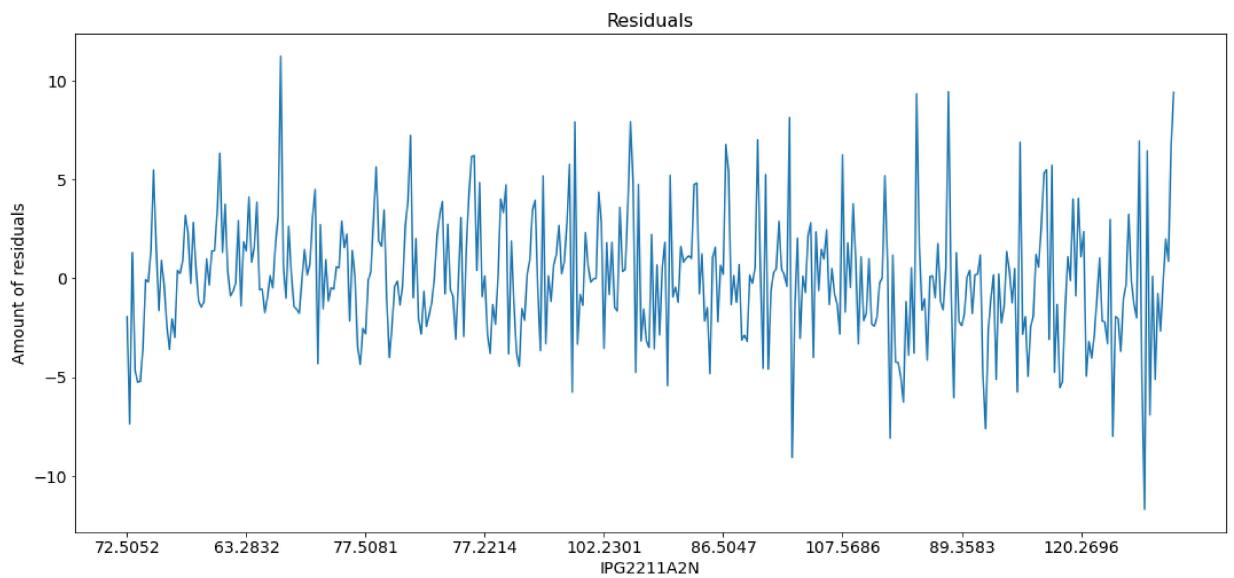
```
In [ ]: As we can see from above models, Model: SARIMAX(2, 1, 2)x(1, 1, 2, 12) fit the best AIC value of 1785 and ma4 has less p-value.
```

```
In [295]: results.plot_diagnostics(figsize=(8,8))
plt.show()
```

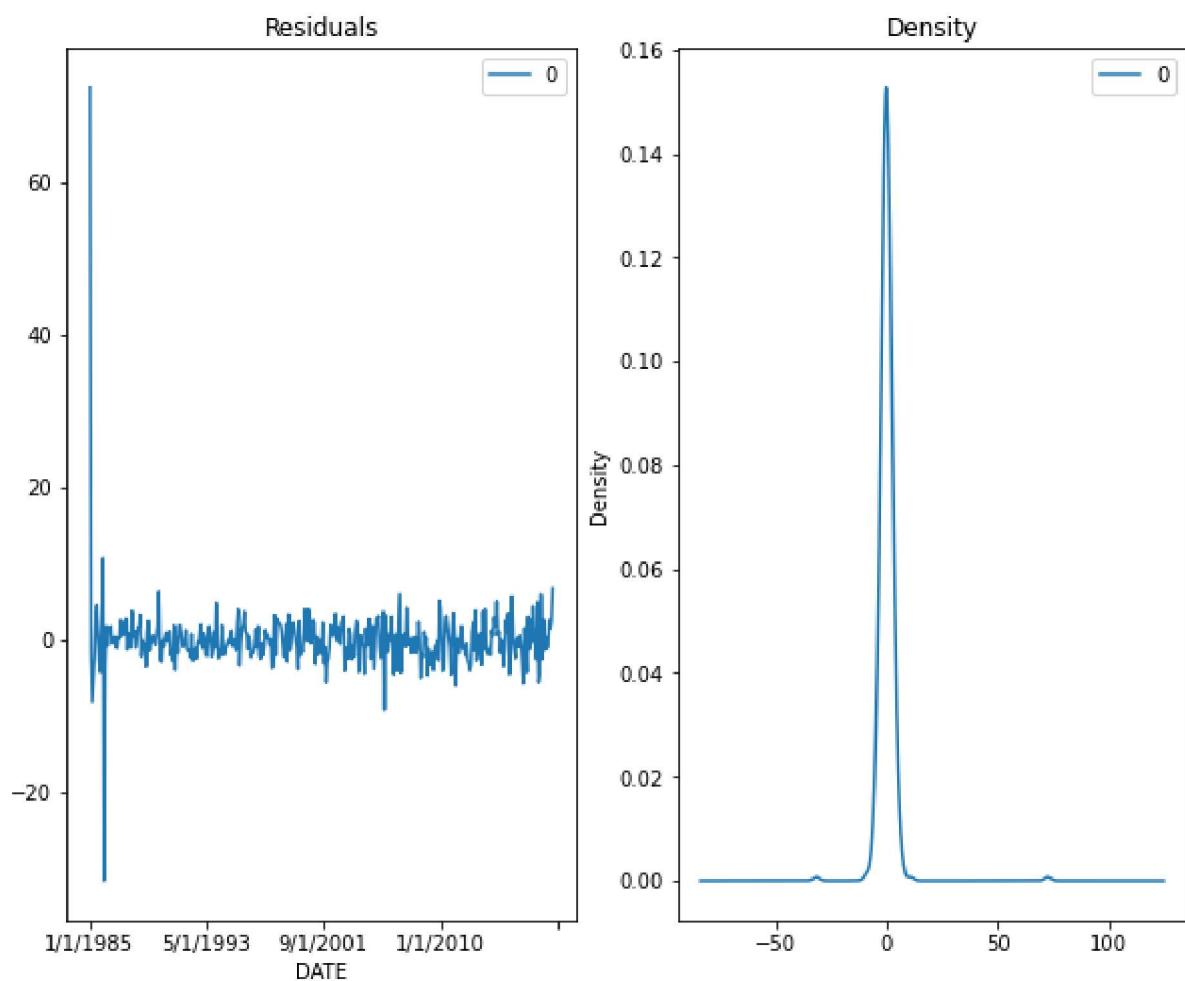


```
In [ ]: All the 4 plots indicates a good fit of the SARIMA model on the given time serie, slightly reduced, which is good. The p-values of the AR1 and MA1 terms have improved next I will plot the residuals to ensure there are no patterns (that is, look for
```

```
In [379]: # Plotting residual
residuals = pd.DataFrame(model2_fit.resid)
plt.figure(figsize=(18, 8))
plt.plot(residuals)
plt.title('Residuals', fontsize=16)
plt.xlabel("IPG2211A2N", fontsize=14)
plt.ylabel("Amount of residuals", fontsize=14)
plt.xticks(np.arange(0, len(data2.IPG2211A2N)+1, 45), labels=[data2.IPG2211A2N[i] for i in range(0, len(data2.IPG2211A2N)+1)])
plt.yticks(fontsize=14)
plt.show()
```



```
In [381]: residuals = pd.DataFrame(model_fit.resid)
fig, ax = plt.subplots(1,2)
residuals.plot(title="Residuals", ax=ax[0])
residuals.plot(kind='kde', title='Density', ax=ax[1])
plt.show()
```



```
In [ ]: The residual errors seem fine with near zero mean and uniform variance
```

```
In [391]: import statsmodels.api as sm
model=sm.tsa.statespace.SARIMAX(train['IPG2211A2N'],order=(4,0,2),seasonal_order=(0,0,0,12))
results=model.fit()
results.summary()
```

Out[391]: SARIMAX Results

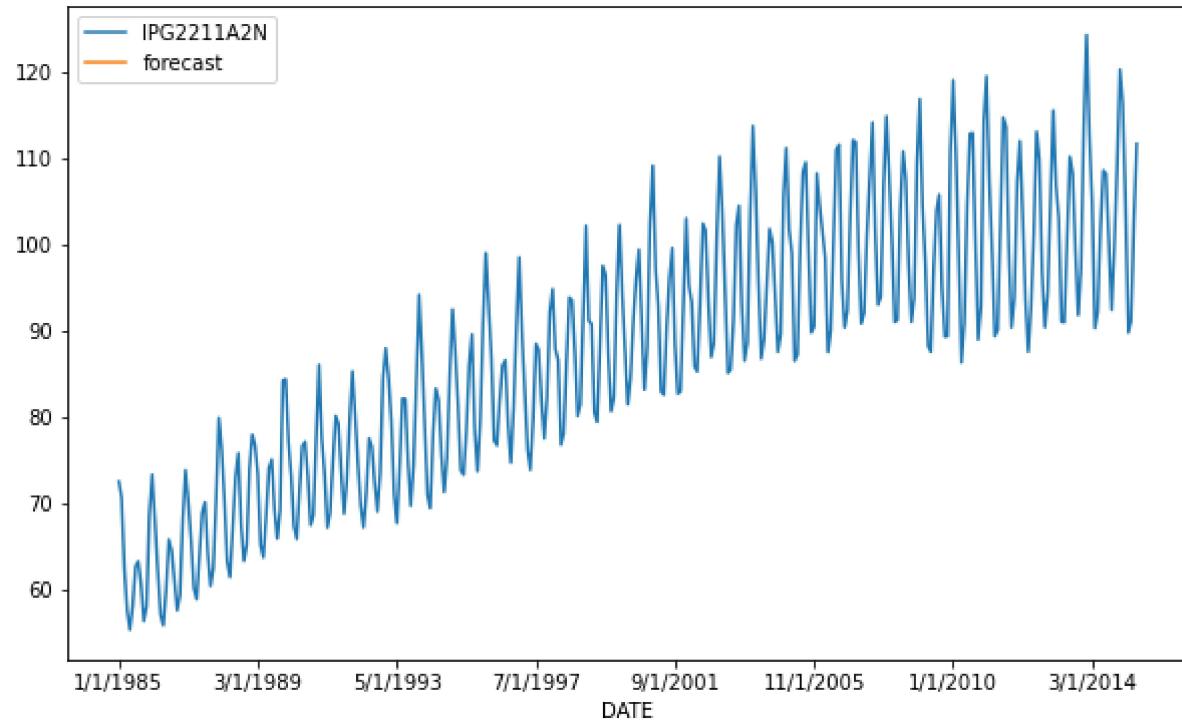
Dep. Variable:		IPG2211A2N		No. Observations:		367
Model:		SARIMAX(4, 0, 2)x(1, 1, [1], 12)		Log Likelihood		-795.452
Date:		Mon, 16 May 2022		AIC		1608.904
Time:		22:58:09		BIC		1643.753
Sample:		01-01-1985		HQIC		1622.768
		- 07-01-2015				
Covariance Type:		opg				
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.5565	0.066	8.442	0.000	0.427	0.686
ar.L2	0.9080	0.084	10.768	0.000	0.743	1.073
ar.L3	-0.5455	0.068	-7.978	0.000	-0.680	-0.412
ar.L4	0.0802	0.057	1.417	0.156	-0.031	0.191
ma.L1	0.0525	0.144	0.364	0.716	-0.230	0.335
ma.L2	-0.9458	0.134	-7.045	0.000	-1.209	-0.683
ar.S.L12	0.0898	0.073	1.235	0.217	-0.053	0.232
ma.S.L12	-0.7655	0.058	-13.118	0.000	-0.880	-0.651
sigma2	5.0095	0.705	7.104	0.000	3.627	6.392
Ljung-Box (L1) (Q):		0.00	Jarque-Bera (JB):	20.28		
Prob(Q):		0.96	Prob(JB):	0.00		
Heteroskedasticity (H):		2.67	Skew:	-0.15		
Prob(H) (two-sided):		0.00	Kurtosis:	4.13		

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [392]: train['forecast']=results.predict(start=300,end=384,dynamic=True)
train[['IPG2211A2N','forecast']].plot(figsize=(10,6))
```

```
Out[392]: <AxesSubplot:xlabel='DATE'>
```



```
In [ ]: Prediction
```

```
In [344]: train= data2.iloc[:30]
test= data2.iloc[-30:]
```

```
In [360]: #I train the model
```

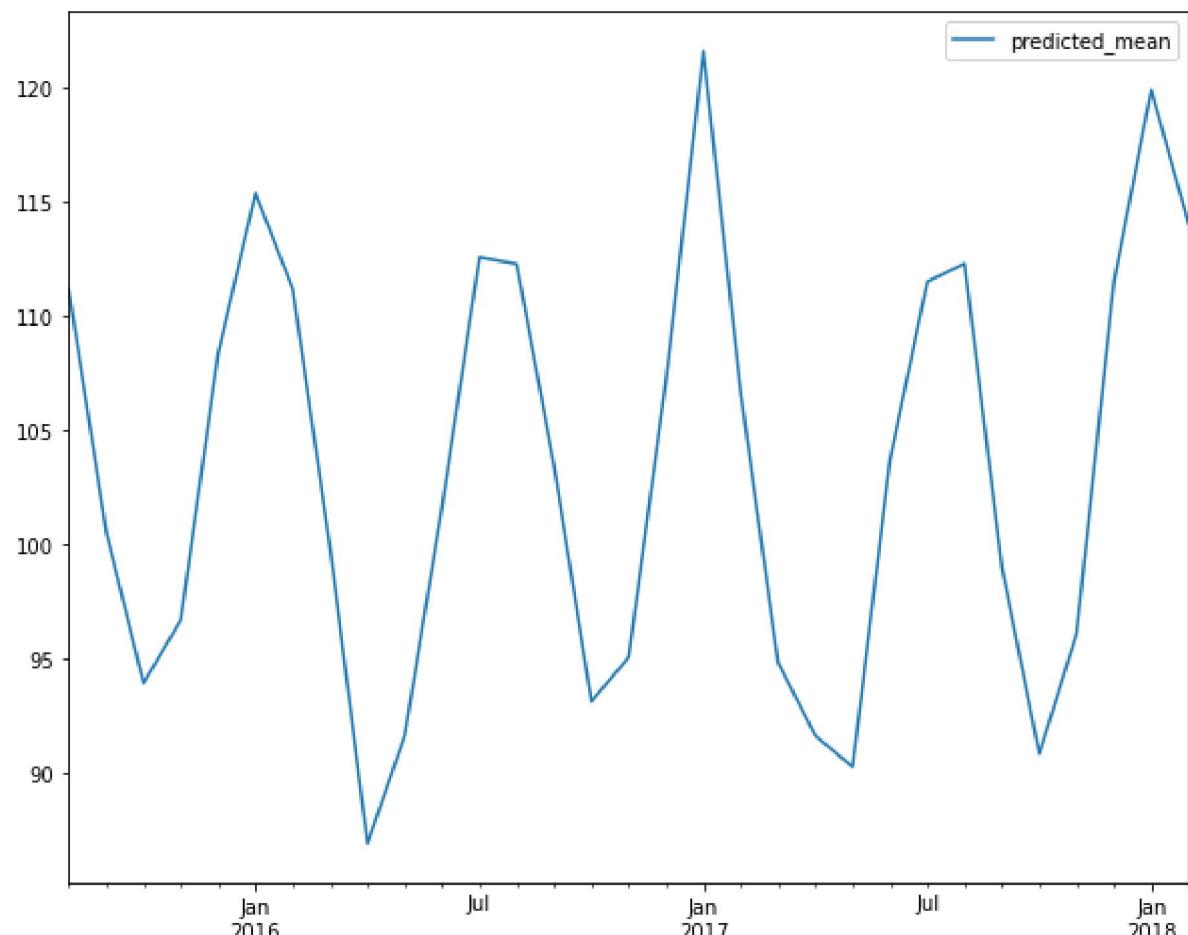
```
model = SARIMAX(data2['IPG2211A2N'],order=(2,1,2), seasonal_order=(1, 1, 2, 12))
results = model.fit()
```

```
In [357]: start=len(train)
end=len(train)+len(test)
predictions = results.predict(start=start, end=end, dynamic=False, typ='levels')
print(predictions)
```

```
2015-08-01    111.158086
2015-09-01    100.557282
2015-10-01     93.913472
2015-11-01     96.688830
2015-12-01    108.370986
2016-01-01    115.382063
2016-02-01    111.180910
2016-03-01    99.835402
2016-04-01    86.886892
2016-05-01    91.621486
2016-06-01    101.650420
2016-07-01    112.576346
2016-08-01    112.281792
2016-09-01    103.445344
2016-10-01    93.117555
2016-11-01    95.053003
2016-12-01    107.220903
2017-01-01    121.601176
2017-02-01    106.679798
2017-03-01    94.822281
2017-04-01    91.614827
2017-05-01    90.246322
2017-06-01    103.711778
2017-07-01    111.509670
2017-08-01    112.300009
2017-09-01    98.990602
2017-10-01    90.816413
2017-11-01    96.092912
2017-12-01    111.491693
2018-01-01    119.919912
2018-02-01    114.031149
Freq: MS, Name: predicted_mean, dtype: float64
```

```
In [358]: predictions.plot(legend=True)
```

```
Out[358]: <AxesSubplot:>
```



```
In [ ]:
```

